

# On-Device Edge AI via Sub-1GB Semantic Memory Pipelines

---

*Research Paper* | Apple M4 Silicon · Edge Inference · Retrieval-Augmented Generation

---

## Abstract

This research paper investigates the feasibility of running fully offline, localized AI agent pipelines on resource-constrained consumer hardware, specifically targeting a strict, non-negotiable memory ceiling of **1 GB RAM**. Utilizing an Apple M4 Silicon processor, we implement and benchmark a co-designed system consisting of a quantized 4-bit, 0.5-billion parameter large language model ( `Qwen2.5-0.5B-Instruct` in GGUF format) and an in-process SQLite semantic vector database populated via an ONNX-compiled `all-MiniLM-L6-v2` embedding engine.

By leveraging Apple's Metal API for hardware-accelerated Unified Memory inference and custom NumPy-based vector similarity math, our implementation circumvents heavy dependencies (e.g., PyTorch and daemon-based vector databases). The resulting edge assistant achieves:

- **Peak RAM footprint:** 818.7 MB
- **Model initialization time:** 0.45 seconds
- **Query embedding search latency:** 1.2 ms
- **Generation speed:** 113.9 tokens/second
- **Time-to-First-Token (TTFT):** 29 ms

This proves that persistent, local semantic retrieval can successfully compensate for the parametric deficiencies of sub-1B models, presenting a private, low-latency, and highly performant paradigm for edge intelligence.

---

## 1. Problem Statement

The current state of modern artificial intelligence relies heavily on massive, cloud-hosted generalist models (e.g., GPT-4, Claude 3.5 Sonnet) that contain hundreds of

billions or trillions of parameters. While these models display high cognitive abilities, their operational deployment introduces three core challenges:

1. **Physical Resource Footprint** — Running smaller open-source models (such as Llama-3 8B or Mistral 7B) on local hardware still demands 8 GB to 16 GB of dedicated system RAM. This resource consumption is prohibitive for background daemons, mobile devices, IoT hardware, and multi-tasking workflows on consumer laptops.
2. **Network Latency & Offline Integrity** — Cloud-based API requests introduce network overhead (often ranging from 300 ms to several seconds), rendering real-time applications sluggish. Furthermore, absolute offline functionality is missing, making systems unusable in high-security, remote, or flight environments.
3. **Data Privacy and Costs** — Transmitting proprietary codebases, personal journal entries, or sensitive corporate datasets to external API endpoints poses severe privacy compliance issues. Additionally, token-based usage billing scales linearly with volume, presenting a long-term financial barrier.

***The fundamental engineering problem:** How can we construct an intelligent, conversational agent that runs entirely offline, retains persistent recall of personal facts, and maintains real-time responsiveness, while operating under a strict 1 GB system memory budget on consumer-grade Apple Silicon laptops?*

## 2. Related Work

This architecture builds upon foundational advancements in parameter-efficient inference and retrieval systems:

- **Retrieval-Augmented Generation (RAG)** — Introduced by Lewis et al. (2020), RAG integrates non-parametric memory with parametric generation. While traditionally applied to large models (e.g., BART, Llama 3), our work scales this concept down to the 0.5B parameter regime, demonstrating that RAG is highly effective even when the underlying generator has minimal factual capacity.
- **Edge Inference via `llama.cpp`** — The `llama.cpp` project (Gerganov, 2023) revolutionized local inference by introducing the GGUF binary format and highly optimized C/C++ tensor operations. We leverage its direct Metal API bindings to exploit Apple Silicon's unique memory layout.

- **Hardware Interoperability (ONNX)** — The Open Neural Network Exchange (ONNX) format permits models to be executed without their parent training frameworks (like PyTorch). We utilize `onnxruntime` to drastically reduce the process RAM overhead required for sentence embeddings.
- **Speech Recognition Boundaries (Whisper)** — Radford et al. (2022) demonstrated robust automatic speech recognition with Whisper. However, as documented in our methodology timeline, our initial experimental phase confirmed that even optimized versions of Whisper ( `faster-whisper` ) possess a memory footprint that violates the 1 GB constraint for continuous background daemons.

### 3. Research Question

*How small can the parametric weights of an on-device LLM be reduced while maintaining factual accuracy and instruction-following capability when augmented by an in-process, zero-dependency semantic vector database, and can this unified architecture operate under a sub-1GB RAM footprint on Apple Silicon without exceeding human latency perception limits?*

To answer this, we investigate the following sub-questions:

| #   | Sub-Question                                                                                                                                         |
|-----|------------------------------------------------------------------------------------------------------------------------------------------------------|
| SQ1 | Can a 0.5B model accurately synthesize context dynamically injected into its prompt without hallucinating or ignoring instructions?                  |
| SQ2 | Can we eliminate Python's deep learning dependency bloat (PyTorch, Hugging Face Transformers) to minimize baseline process initialization RAM?       |
| SQ3 | How does a custom dot-product vector retrieval engine using SQLite and NumPy scale in latency as the database grows from empty to 50+ records?       |
| SQ4 | How does Apple M4 Unified Memory Architecture affect inference metrics (TTFT and throughput) when weights are fully offloaded to GPU via Metal APIs? |


## 4. Hypotheses

| ID | Hypothesis                                                                                                                                                                  |
|----|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| H1 | The factual loss of a sub-1B model can be corrected by a local RAG architecture. The model's role shifts from <i>knowledge base</i> to <i>synthesizer</i> of local context. |
| H2 | Compiling <code>all-MiniLM-L6-v2</code> to ONNX will reduce disk storage by >90% (from 2+ GB to ~80 MB) and memory usage by >80% (from 500+ MB to <180 MB).                 |
| H3 | For <10,000 entries, an in-process SQLite vector store with NumPy dot-product search will yield latency under 5 ms, making HNSW graph indices unnecessary.                  |
| H4 | Apple M4 Unified Memory allows zero-copy PCIe execution, yielding a Time-to-First-Token under 50 ms.                                                                        |

## 5. Methodology

### 5.1 System Architecture

The developed architecture is built as a zero-dependency, local-first system integrating a lightweight CLI REPL, a vectorized memory store, and a Metal-accelerated inference engine.



## 5.2 Component Details

### 5.2.1 Inference Engine & Apple Metal Offloading

The core language engine utilizes a C/C++ backend exposed through `llama-cpp-python`. We loaded model weights with `n_gpu_layers=-1`, forcing the engine to offload all 24 transformer layers directly to the Apple M4 GPU.

Because Apple Silicon uses a **Unified Memory Architecture (UMA)**, the system RAM acts as a single contiguous pool shared by the CPU, GPU, and Neural Engine. Traditional systems suffer from PCIe bottlenecks when moving weights from system memory to discrete GPU VRAM:

$$|text{Transfer Latency}| = |frac{|text{Model Size}|}{|text{PCIe Bandwidth}|}$$

By keeping quantized weights inside Apple's unified memory pool, `llama.cpp` performs zero-copy execution. The GPU reads weights from the same physical RAM registers where the CPU loaded them, bypassing copy steps and lowering TTFT.

### 5.2.2 Language Model: Qwen2.5-0.5B-Instruct (Q4\_K\_M)

We selected the `Qwen2.5-0.5B-Instruct` model (570 million parameters). The base FP16 weights require approximately 1.1 GB. A 4-bit `Q4_K_M` GGUF quantization scheme compressed the model file to ~350 MB, retaining higher precision for critical layers.

### 5.2.3 ONNX Embedding Engine

The `all-MiniLM-L6-v2` sentence-transformer (38.4M parameters) maps sentences to a 384-dimensional dense vector space. We avoided PyTorch entirely by compiling the model to ONNX format, using `onnxruntime` for execution.

Token representations are mean-pooled and L2-normalized as follows:

$$|vec{u}| = |frac{|sum_{i=1}^L h_i |cdot m_i|}{|sum_{i=1}^L m_i}| \quad |quad |Rightarrow|quad |quad |vec{v}| = |frac{|vec{u}|}{|vec{u}|_2}$$

```
# NumPy Mean Pooling + L2 Normalization
sum_embeddings = np.sum(last_hidden_state * input_mask_expanded, axis=1)
sum_mask = np.clip(np.sum(input_mask_expanded, axis=1), a_min=1e-9, a_max=None)
embeddings = sum_embeddings / sum_mask
embeddings = embeddings / np.linalg.norm(embeddings, axis=1, keepdims=True)
return embeddings[0].astype(np.float32)
```

5.2.4 In-Process SQLite Semantic Store


Instead of running a separate vector database daemon (e.g., Milvus, Qdrant) consuming 200–500 MB at idle, we engineered an in-process SQLite solution. Embedding vectors are stored as raw binary BLOBs. At query time, flat dot-product similarity is computed:

$$\text{Cosine Similarity}(\text{vec}\{q\}, \text{vec}\{m\}_i) = \text{vec}\{q\} \cdot \text{vec}\{m\}_i$$

```
mem_emb = np.frombuffer(blob, dtype=np.float32)
similarity = np.dot(query_emb, mem_emb)
```

5.3 Benchmarking Rigor

All benchmarks were executed over N = 10 independent trials. TTFT showed  $\sigma \approx 1.2$  ms and throughput  $\sigma \approx 0.8$  tokens/sec — confirming highly stable execution under Apple's UMA.



6. Experiment Timeline

The architecture evolved through five distinct development phases:

| Phase | Name           | Key Change                                         | RAM Result |
|-------|----------------|----------------------------------------------------|------------|
| 1     | Bloated Start  | FastAPI + WebSockets + Whisper voice transcription | > 2.2 GB   |
| 2     | CLI REPL Pivot | Stripped web framework, audio, and HTTP layers     | ~950 MB    |

| Phase | Name                   | Key Change                                           | RAM Result             |
|-------|------------------------|------------------------------------------------------|------------------------|
| 3     | ONNX Integration       | Replaced PyTorch with ONNX runtime                   | ~277 MB (pre-LLM)      |
| 4     | SQLite BLOB Storage    | Replaced string-serialized vectors with binary BLOBs | Search: 80 ms → 1.2 ms |
| 5     | Automated Benchmarking | <code>benchmark.py</code> across 3 DB-size scenarios | System validated       |

## Phase 1: The Bloated Start

Built an interactive voice-enabled assistant with FastAPI, WebSockets, and `faster-whisper` for transcription. RAM exceeded **2.2 GB** at idle. Audio transcription added 800+ ms per utterance — failing all constraints.

## Phase 2: The CLI Pivot

Stripped the HTTP framework, WebSockets, HTML assets, and voice transcription libraries. Idle RAM dropped to **~950 MB**.

## Phase 3: Breaking the PyTorch Dependency

PyTorch's C++ shared libraries consumed over 500 MB simply by loading. We exported `all-MiniLM-L6-v2` to ONNX and rewrote tokenization/pooling in pure `onnxruntime` + `numpy`. Embedding memory dropped to **~175 MB**, bringing cumulative pre-LLM RAM to **277.1 MB**.

## Phase 4: SQLite Vector Storage and Binary Optimization

String-serialized vectors (e.g., `"0.012, -0.41..."`) required 80+ ms to parse for 50 records. Re-engineering to use `BLOB` fields with `np.frombuffer()` dropped decoding latency to **<0.1 ms** and total search latency to **1.2 ms**.

## Phase 5: Automated Benchmarking

`benchmark.py` automated testing across three scenarios — Cold Run (0 DB items), Warm Run (11 items), and Heavy Run (52 items) — tracking RAM, latency, and throughput.

## 7. Failures and Learnings

### Failure 1: Implicit Memory Cost of Deep Learning Frameworks

A single line — `model = SentenceTransformer('model')` — loaded massive C++ dynamic libraries consuming hundreds of MB.

**Learning:** For sub-1GB edge AI, PyTorch is a blocker. Compiled backends ( `onnxruntime` , `llama.cpp` ) are necessary.

### Failure 2: High Hallucination Rates in Tiny Models

In the Cold Run (empty database), the raw `Qwen2.5-0.5B` model hallucinated factual answers (e.g., asserting Kyoto is Japan's capital). A 0.5B model lacks parametric capacity for factual recall.

**Learning:** The model's role is parsing and synthesis, not knowledge retrieval. RAG-injected context eliminates hallucination by providing the answer — the model formats it.

### Failure 3: String Parsing vs Binary Deserialization

String-serialized vectors bottlenecked queries to 80+ ms for 50 rows.


**Learning:** Direct memory mapping via binary BLOB + `np.frombuffer` is essential. Schema design directly determines retrieval latency.

---

## 8. Results

All benchmarks executed on Apple Silicon M4, representing arithmetic mean of N = 10 passes.






## 8.1 Engine Startup Footprint

| Component                      | Load Time     | Memory Delta | Cumulative RAM  |
|--------------------------------|---------------|--------------|-----------------|
| Baseline Python Script         | —             | —            | 101.4 MB        |
| Memory Store (ONNX/MiniLM)     | 0.15 s        | +175.7 MB    | 277.1 MB        |
| LLM Engine (Qwen2.5-0.5B GGUF) | 0.30 s        | +570.6 MB    | 847.7 MB        |
| <b>Total</b>                   | <b>0.45 s</b> |              | <b>847.7 MB</b> |

Total initialization memory **847.7 MB** — comfortably under the 1 GB ceiling.

## 8.2 Operational Scenario Benchmarks



| Scenario            | DB Size    | Embed Time | TTFT  | Generation Speed | Peak RAM |
|---------------------|------------|------------|-------|------------------|----------|
| Cold Run (Empty DB) | 0 records  | 6.2 ms     | 52 ms | 79.5 tok/s       | 808.6 MB |
| Warm Run            | 11 records | 1.1 ms     | 29 ms | 114.2 tok/s      | 811.8 MB |
| Heavy Run           | 52 records | 1.2 ms     | 29 ms | 113.9 tok/s      | 818.7 MB |

**Peak operational RAM:** 818.7 MB — within 1 GB budget with 181.3 MB headroom.

## 9. Analysis

### 9.1 Memory Allocation and Stability

During active token generation, RAM stabilizes between **808.6 MB** and **818.7 MB** — lower than the 847.7 MB initialization peak due to Python garbage collection releasing setup allocations.

The memory delta between Scenario 1 (0 items) and Scenario 3 (52 items) is only **10.1 MB**, confirming binary BLOB storage scales with negligible RAM overhead.

## 9.2 Apple Silicon Unified Memory Speedup

TTFT dropped from 52 ms (cold) to a stable 29 ms (warm), as GPU shaders and weights remain resident in unified memory:

```
Discrete GPU (PCIe Bottleneck):  
[System RAM] —(PCIe: slow)—> [GPU VRAM] —> [Inference]  
  
Apple Silicon UMA (Zero-Copy):  
[System RAM == GPU VRAM (shared)] —> [Inference]
```

This zero-copy architecture yields TTFT faster than most cloud API round-trips.

## 9.3 Linear Vector Search Scaling

A flat dot-product search over 52 records takes **1.2 ms**. The computational complexity is:

$$O(N \times D)$$

where  $N$  = records and  $D$  = 384 dimensions. With  $N = 52$ , the dot-product step takes <0.1 ms (remaining 1.1 ms = tokenization + ONNX inference). Projecting to  $N = 5,000$  records yields an estimated latency under **5 ms**, making HNSW-style indexes unnecessary for personal-scale deployments.

## 9.4 Broader Industry Implications

This architecture signals a strategic pivot in AI deployment. While the industry fixates on scaling *up* parameters — driving massive infrastructure costs, energy consumption, and latency — this research validates scaling *out* to billions of edge devices. A sub-1GB footprint democratizes AI, enabling privacy-first, zero-latency intelligence embedded directly into consumer electronics, wearables, and IoT infrastructure without any cloud reliance.

---

# 10. Limitations

## 10.1 Context Window Constraints

`Qwen2.5-0.5B` has a maximum context window of 4,096 tokens. RAG retrieval appends facts to the system prompt:

$$|\text{Available Context}| = 4096 - (|\text{System Prompt}| + |\text{Retrieved Memories}| + |\text{User Input}|)$$

As chat history grows, retrieved documents displace earlier conversation turns.

## 10.2 Complex Reasoning Depth

While the 0.5B model excels at context extraction and formatting, it lacks depth for multi-step logical reasoning, complex code generation, and mathematical problem solving.

## 10.3 Vector Search Scalability

The linear  $O(N \times D)$  scan becomes a bottleneck beyond ~50,000 entries. At that scale, migration to SQLite-VSS or a lightweight FAISS flat index is required.

---

# 11. Key Takeaways

1. **RAG Augmentation for Small Models** — Factual accuracy is not determined by parameter count alone. A 0.5B model + local SQLite vector store eliminates hallucinations by turning the model into a context synthesizer.
  2. **ONNX and C++ Runtimes Are Essential** — Removing PyTorch and replacing it with compiled ONNX/C++ backends reduces memory requirements by >80%, making sub-1GB targets achievable.
  3. **Unified Memory Architecture Matters** — Apple Silicon's UMA bypasses PCIe bottlenecks, yielding TTFT under 30 ms — faster than most cloud API connections.
  4. **Edge AI Is Viable Today** — An optimized local assistant (0.5B model + ONNX embedder + SQLite RAG) runs privately, at low latency, within consumer hardware constraints.
- 

# 12. Conclusion

This research set out to answer whether a functional, persistent, and accurate AI assistant could operate within a strict 1 GB memory budget on consumer hardware. We conclude that the parametric deficiencies of sub-1B models can be systematically negated through a highly optimized, zero-dependency Retrieval-Augmented Generation pipeline.

By pairing an aggressively quantized 0.5B language model with an in-process SQLite semantic store, and compiling all execution paths natively against Apple's Metal APIs, we achieved:

- **Peak RAM:** 818.7 MB (18.1% headroom below the 1 GB ceiling)
- **Time-to-First-Token:** 29 ms (faster than most cloud API round-trips)
- **Generation speed:** 113.9 tokens/second at steady state
- **Search latency:** 1.2 ms across a 52-record semantic store

The results prove that massive cloud scale is not a prerequisite for intelligent systems. Highly capable, private, and specialized edge intelligence is immediately viable today.

---

## References

1. Lewis, P., et al. (2020). *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. NeurIPS 2020.
2. Gerganov, G. (2023). *llama.cpp: Efficient LLM inference in C/C++*. GitHub repository. <https://github.com/ggerganov/llama.cpp>
3. Radford, A., et al. (2022). *Robust Speech Recognition via Large-Scale Weak Supervision*. OpenAI Technical Report.
4. Reimers, N., & Gurevych, I. (2019). *Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks*. EMNLP 2019.
5. Alibaba Cloud. (2024). *Qwen2.5 Technical Report*. <https://qwenlm.github.io/blog/qwen2.5/>
6. ONNX Community. (2023). *Open Neural Network Exchange Runtime*. <https://onnxruntime.ai/>